


CS323 Operating Systems


Operating Systems Summary

Mathias Payer and Sanidhya Kashyap

EPFL, Fall 2021



Virtualization: Summary




CPU Virtualization: Processes and Scheduling

- Processes are a purely virtual concept
- Separating policies and mechanisms enables modularity
- Schedulers need to optimize for different metrics: utilization, turnaround, response time, fairness and forward progress
 - FIFO: simple, non-preemptive scheduler
 - SJF: non-preemptive, prevents process jams
 - STFC: preemptive, prevents jams of late processes
 - RR: preemptive, great response time, bad turnaround
 - MLFQ: preemptive, most realistic
 - CFS: fair scheduler by virtualizing time
- Past behavior is good predictor for future behavior

Memory Virtualization: Segmentation and Paging

- OS manages access to constrained resources
 - Principle: limited direct execution (bare metal when possible, intercept when needed)
 - **CPU**: time sharing between processes (low switching cost)
 - **Memory**: space sharing (disk I/O is slow, so time sharing is expensive)
- **Fragmentation**: space lost due to internal or external padding
- **Paging**: MMU fully translates between virtual and physical addresses
 - One flat page table (array)
 - Multi-level page table
 - Pros? Cons? What are size requirements?
- Paging and swapping allows process to execute with only the working set resident in memory, remaining pages can be stored on disk




- Abstraction: locks to protect shared data structures
- Mechanism: interrupt-based locks
- Mechanism: atomic hardware locks
- Busy waiting (spin locks) versus wait queues
- Condition variables
- Semaphores
- Signaling through condition variables and semaphores

Difference parallelism and concurrency

- **Parallelism:** multiple threads (or processes) working on a single task using multiple CPU cores (i.e., stuff happens at the same physical time)
- **Concurrency:** tasks can start, run, and complete in overlapping time periods (i.e., tasks run at the same virtual time)

Concurrency summary


- Spin lock, CV, and semaphore synchronize multiple threads/processes
 - **Spin lock:** atomic access, no ordering, spinning
 - **Condition variable:** atomic access, queue, OS primitive
 - **Semaphore:** shared access to critical section with (int) state
- All three primitives are equally powerful
 - Each primitive can be used to implement both other primitives
 - Performance may differ!
- Synchronization is challenging and may introduce different types of bugs such as atomicity violation, order violation, or deadlocks.



- Device interaction and device drivers
- IO Scheduling and harddrive throughput
 - Disk layout
 - Disk virtualization (RAID)
- Filesystem API
- Filesystem implementation
 - Inodes and devices
 - File descriptors
 - File names
- Crash resistance
- Journaling

- Overlap IO and computation as much as possible!
 - Use interrupts
 - Use DMA
- Driver classes provide common interface
- Storage: read/write/seek of blocks
- Minimize random IO (i.e., quick sort is really bad on HDDs)
- Carefully schedule IO on slow devices
- RAID virtualizes disks

- Filesystem API: handle interaction with the file system
- Three ways to identify a file
 - File names (for humans)
 - Inodes and devices (on the disk)
 - File descriptors (for a process)
- Filesystem implementation
 - Inodes for metadata
 - Bitmaps for inodes/data blocks
 - Superblock for global metadata
- Crash resistance: filesystem check (FSCK)
- Journaling: keep track of metadata, enforce atomicity
 - All modern filesystems use journaling
 - FSCK still useful due to bitflips/bugs



Two topics: testing and mitigations

- Testing helps developers find as many bugs as possible
 - Fuzzing generates test cases
 - Sanitization detects policy violations
- Mitigations detect policy violations at runtime, stop exploits

- Software testing finds bugs before an attacker can exploit them
- Manual testing: write test cases to trigger exceptions
- Fuzz testing automates and randomizes testing
- Sanitizers allow early bug detection, not just on exceptions
- AddressSanitizer is the most commonly used sanitizer and enforces probabilistic memory safety by recording metadata for every allocated object and checking every memory read/write.

- Deployed mitigations do not stop all attacks
- ***Data Execution Prevention*** stops code injection attacks, but does not stop code reuse attacks
- ***Address Space Layout Randomization*** is probabilistic, shuffles memory space, prone to information leaks
- ***Stack Canaries*** are probabilistic, do not protect against direct overwrites, prone to information leaks
- ***Control-Flow Integrity*** restricts control-flow hijack attacks, does not protect against data-only attacks



Figure 1: Understanding OS' will help your career!

- Learn core concepts
- Become aware of design decisions and policies
 - Virtualization: CPU and Memory
 - Concurrency: performance trade-offs
 - Persistence: correctness and recovery
 - Security: software testing versus mitigations

Learning goals: labs

- Lab 0: practice C programming and debugging
- Lab 1: thread scheduling and memory allocation
- Lab 2: concurrency and message passing
- Lab 3: simple file system
- Lab 4: software security testing

But the main goal was to become better programmers, i.e., using a specification to implement and test a prototype, then integrate it into the overall system.

- Open book, online, 3 hours, Jan29 8:15 to 11:15

- Open book, online, 3 hours, Jan29 8:15 to 11:15
- Several questions for each topic
 - Theory: based on the lectures
 - Practice: based on the labs
 - 2-4 questions per topic (between 8 and 16 questions)

- Open book, online, 3 hours, Jan29 8:15 to 11:15
- Several questions for each topic
 - Theory: based on the lectures
 - Practice: based on the labs
 - 2-4 questions per topic (between 8 and 16 questions)
- Answers will be
 - Numbers (e.g., 15 pages)
 - Code (merging for buddy allocator)
 - Prose (why you don't temporally separate memory)

How to cheat

- Build a community/cluster
- Share a communication channel
 - Google Docs is ideal for collaborative editing
 - Shared voice also helps (team speak)
- Pay someone else to solve the exam for you

How to mitigate cheating

Goal: limit cheating while ***minimizing disruption*** for honest students.

- Rely on honor system

How to mitigate cheating

Goal: limit cheating while ***minimizing disruption*** for honest students.

- Rely on honor system
- Make it cost ineffective to cheat

How to mitigate cheating

Goal: limit cheating while ***minimizing disruption*** for honest students.

- Rely on honor system
- Make it cost ineffective to cheat
- Make sharing hard
 - You get one of many different exam versions
 - Number/code questions: subtle changes to question
 - Prose questions: anti-plagiarism detection

How to mitigate cheating

Goal: limit cheating while ***minimizing disruption*** for honest students.

- Rely on honor system
- Make it cost ineffective to cheat
- Make sharing hard
 - You get one of many different exam versions
 - Number/code questions: subtle changes to question
 - Prose questions: anti-plagiarism detection
- Prohibit imposters
 - Arbitrary oral exams will replace written exams
 - Random selection of students
 - Suspicious cases

- You will get link to read-only Google Doc on Moodle
- Login with your EPFL account
- Make a private copy, solve exam in Google Doc
- Submit PDF of your Google Doc through Moodle
- If your internet dies
 - Submit as screenshots via phone
 - Give us access to your document

Final Exam: Announcements

- Zoom for announcements and questions
 - Voice is easier for your questions
 - Higher throughput
- We will monitor slack as well


Prepare for the exam

- Watch the lectures and read the book chapters
- Solve the exercises
- Solve the practice midterm/final
- Create summaries and indexes

Prepare for the exam

- Watch the lectures and read the book chapters
- Solve the exercises
- Solve the practice midterm/final
- Create summaries and indexes
- We offer office hours to answer questions
 - January, each Wed 10-11 on zoom
 - No exercises this Wednesday

All done?



- Feedback is appreciated, be as detailed as possible
 - For good statistics, I need all of you to respond!
 - Be open and positive!
- What was great? What can be improved?
- Be as detailed as possible
 - How can I improve the class?
 - How can I improve the labs?
 - Was the workload and distribution reasonable?
- Let me know what else you were missing!

All done?



Figure 3: Keep your curiosity going, the HexHive lab offers fun {BSc|MSc} software/systems security projects. Talk to us!